

The Interface Is a Variable: Measuring the Cost and Reliability of Purpose-Built UI Representations for LLM Agents

Abstract

Language-model agents usually operate software through interfaces built for people. They read screenshots, accessibility trees, or raw HTML, and they act by clicking and typing. This paper treats the interface representation, the form in which an application shows its state and its available actions to an agent, as an experimental variable rather than as a fixed cost of automation. We hold a small online store, its ten tasks, and a deterministic grader constant, and present the same store to the same model in four ways: a screenshot (C1), an accessibility tree (C2), flat tools (C3), and a purpose-built JSON view document (C4) that states which screen the agent is on and which actions are allowed right now. We measure task success, token cost, step count, and illegal actions, and we read the results as a cost-reliability frontier: a representation earns its place only if nothing else is both cheaper and at least as reliable. On the first model measured, the two human surfaces are dominated. The screenshot (C1) costs roughly ten times the tokens of the cheapest condition and completes fewer than a third of its tasks. The accessibility tree (C2) is cheaper than pixels but error-prone. Between the structured conditions the verdict is closer: the view document (C4) finishes tasks in fewer steps, but it spends more tokens than flat tools (C3) and does not improve success, so flat tools (C3) is preferred on this store. All numbers come from one general model (`gemini-2.5-flash`) on one synthetic store. We present them as a measurement method and a first data point, and we identify the conditions under which a purpose-built agent surface should begin to pay off.

1. Introduction

An agent that operates software runs a loop. It looks at the application, decides on one action, performs that action, and looks again. Every look enters the model as input tokens, and every decision is billed as output tokens. Across the loop the application itself does not change. The only thing that changes is the form in which the agent sees it.

That form is almost always inherited from software built for people. Three channels dominate practice today, and each fails the agent in its own way.

1. **Screenshot (C1).** The model sees pixels and must translate intent into click coordinates. Grounding those clicks is a well-documented failure mode of GUI agents (SeeClick, UI-TARS, Claude computer use). Worse, image cost tracks the size of the viewport rather than

the amount of information on the page, so a blank page and a busy page at the same window size cost exactly the same.

2. **Accessibility tree (C2).** The model reads the text tree that assistive technology uses. The tree is verbose. Mind2Web filters it because it is too large to send whole, and UIFormer measures interface text at 80 to 99 percent of an agent's total token cost. The tree also names controls that the backend will reject, which turns confident clicks into errors.
3. **Flat tools (C3).** The model receives a catalog of function schemas and a short status object after each call, with no description of the current screen. This is close to how the Model Context Protocol exposes capabilities. It is cheap per call once the schema is paid for, but the model must remember on its own which view it is on, which actions are currently valid, and which arguments are legal.

None of these channels was designed to answer, cheaply and unambiguously, the two questions an agent asks at every step: *what is true right now, and what may I do next?*

The alternative we measure is a **view document (C4)**: a JSON object generated from backend state that names the current view, lists the entities on it, and lists every action with an `enabled` flag and a schema for its arguments. It is not a better screenshot, and it is not a pruned DOM. It is a projection of the same backend state that the grader already trusts. Section 4.1 shows a real one.

Whether applications should ship such a surface, the way they already ship an API, is the motivating product question. It is not the contribution. The contribution is a measurement. We hold the store, the tasks, the grader, and the model fixed, vary only the representation, and read the outcome on a cost-reliability frontier, where a representation survives only if nothing else is both cheaper and at least as reliable.

The experiment in one paragraph. One synthetic store (MiniShop), ten frozen tasks (seven that should succeed, three that should be refused), one deterministic grader that inspects backend order records, one model at a time, temperature zero, at most twenty steps per task. The agent's memory is the task text, its prior actions, and the current observation. The only thing that varies is how the agent sees and acts: screenshot (C1), accessibility tree (C2), flat tools (C3), or view document (C4). The two structured conditions expose exactly the same seven operations, so any difference between them is the document, not a different set of buttons.

What we found, in brief. On the first model measured, the human surfaces lose decisively. The screenshot (C1) never completes a purchase task and costs about 46,000 input tokens per task. The accessibility tree (C2) succeeds more often but commits the most illegal and malformed actions. Flat tools (C3) completes every task. The view document (C4) nearly matches it, uses the fewest steps of any condition, costs more tokens than flat tools (C3), and fails exactly one task, because the model ignored an action that the document had clearly marked as unavailable. These numbers are one cell of a larger planned grid, one model crossed with one application, and we treat them accordingly.

Contributions.

1. A controlled method for comparing interface representations at *matched action grain*: flat tools (C3) and view document (C4) expose identical operations, so a win for the document cannot be explained by a smaller or simpler action set.
2. A cost-reliability reading with paired per-task inference, rather than a success-only leaderboard: conditions are compared task by task with bootstrap confidence intervals.
3. First evidence on one store and one model, anchored by a model-free observation-cost baseline (what each representation costs to read before any model is involved) and a deterministic execution grader.

Research questions.

- **RQ1.** Does the representation change token cost, step count, task success, and illegal-action rate on a fixed task set?
- **RQ2.** Is some representation cheaper and at least as reliable as reusing the human page? In particular, does the view document (C4) dominate the screenshot (C1) and the accessibility tree (C2)?
- **RQ3.** Does the ordering of conditions hold across general models, or is it an artifact of one model?

The matching hypotheses: **H1**, cost and reliability differ across conditions. **H2**, the view document (C4) sits on the cost-reliability frontier and the screenshot (C1) is dominated. **H3**, the ordering of conditions is stable across general models even if the size of the gaps is not. On this first model, H1 holds at a coarse grain (human surfaces versus structured ones). H2 does not hold in full: the view document (C4) beats the human surfaces but not flat tools (C3). H3 is untested.

How the paper is organized. Section 2 collects the terms used throughout. Section 3 places the study among prior work. Section 4 describes the store, the conditions, the tasks, and a real view document (C4). Section 5 defines the metrics and the analysis. Section 6 reports results in three layers: a model-free baseline, an oracle control, and the model run. Section 7 separates what this cell shows from what later cells must show. Section 8 describes how a real application could ship a view document (C4). Sections 9 and 10 cover limitations and future work.

2. Terms used in this paper

Conditions are always written as method name plus identifier.

- **Screenshot (C1).** A rendered image of the human page. The agent clicks coordinates, types, and scrolls.
- **Accessibility tree (C2).** The text tree of that page (roles and names). The agent clicks or fills named elements.

- **Flat tools (C3).** A catalog of callable functions plus a short status object per call. No description of the current view.
- **View document (C4).** A JSON object with `view`, `state`, `entities`, and `affordances`, generated from backend state.
- **Affordance.** One action the interface exposes, such as `open_product` or `pay`, with an `enabled` flag and an argument schema.
- **Matched action grain.** Flat tools (C3) and view document (C4) expose the same seven operations, so differences between them come from the document alone.
- **Illegal action.** An action the backend rejects.
- **Malformed action.** A model reply that is not a usable action at all (wrong shape). Counted separately from illegal actions.
- **Ignored affordance.** The model attempts an action the document marked `enabled: false`. Task t10 is the running example.
- **Step.** One observe, decide, act cycle. Runs are capped at 20 steps.
- **Execution-based grader.** A deterministic check over backend order records. Never another model.
- **Cost-reliability frontier.** Token cost plotted against success. A condition is *dominated* if another condition is no more expensive and at least as reliable.
- **Refusal task.** A task whose correct outcome is to decline (for example, the requested size is sold out). Reported separately, because doing nothing can pass one.

3. Related work

Most prior work asks which *agent* is better at operating a computer and treats the interface as a fixed property of the benchmark. This study inverts that. The model is held fixed and the interface is the variable.

Reusing the human page. DOM and accessibility-tree agents (Mind2Web, WebArena) correspond to our accessibility tree (C2). Screenshot and pixel systems (Claude computer use, UI-TARS, SeeClick) correspond to our screenshot (C1). In those lines of work, grounding is the central difficulty, and screenshots cost on the order of one to two thousand tokens each. Our question is not how to ground pixels better. It is whether a non-pixel surface, authored for the agent, is better for ordinary general models. Specialized pixel-native models are deferred to a later check.

Compressing the human tree. UIFormer and Prune4Web prune and restructure the human DOM or accessibility tree to save tokens. UIFormer reports interface text at 80 to 99 percent of agent token cost and reductions around 49 to 56 percent. This is the closest neighboring line, and it improves the accessibility tree (C2). The view document (C4) differs in origin: it is generated from backend state, not distilled from the human page, and we measure it as a controlled experimental arm rather than as a plug-in optimizer.

Agent-facing protocols. A2UI runs in the opposite direction, letting an agent generate interface for a human to view. The Model Context Protocol standardizes tool catalogs, which is close to flat tools (C3). The view document (C4) is application-to-agent, per view, and carries validity constraints, and we match it in grain to flat tools (C3) so the document's added value can be isolated. We use it as a measurement probe, not as a proposed standard.

Benchmarks. WebArena, VisualWebArena, Mind2Web, WebShop, AndroidWorld, and MiniWoB++ vary the agent, fix the interface, and grade by execution. We reuse the execution-grading idea and invert the design. We also report a model-free observation-cost baseline, which capability benchmarks do not provide.

Table R1. Positioning against each strand (verified sources only).

Strand	Representative work	What it does	How this study differs
DOM / accessibility agents	Mind2Web; WebArena	Operate the human page from HTML or a tree	That channel is our accessibility tree (C2), one of four controlled conditions
Screenshot / pixel computer use	Claude computer use; UI-TARS; SeeClick	Perceive pixels, emit coordinates	That channel is our screenshot (C1). We study representation, not better grounding
Optimizing the human tree	UIFormer; Prune4Web	Compress a human-derived tree	Closest neighbor. The view document (C4) is generated from backend state instead
Agent-UI protocols	A2UI; MCP	A2UI is agent to human; MCP is a tool catalog	The view document (C4) is application to agent and matched in grain to flat tools (C3)
GUI / web benchmarks	WebArena; VisualWebArena; Mind2Web; WebShop; AndroidWorld; MiniWoB++	Vary the agent, fix the interface	We fix the agent and vary the interface

What this study does not claim. That representation matters is already known: Mind2Web filters the DOM because it is too large, SeeClick is motivated by lengthy HTML, and UIFormer measures how much of an agent's budget the interface consumes. Execution grading is not new. Token-efficient interface trees are not new. Agent-facing surfaces are already discussed (A2UI, MCP). Masking invalid actions is a known idea in reinforcement learning and tool use, and we do not attach a citation we have not verified. The claim here is the controlled four-way comparison on a fixed application, model, and grader, and the scope of that claim is bounded by the setup: one store, ten tasks, one model, and a hand-authored view document (C4).

4. Study design

The store. MiniShop is a small shop with a catalog, product pages, a cart, and a checkout. Its backend state is the single source of truth for grading. The human pages carry realistic clutter (promotional copy, a search box, decorative filters, help links), so the screenshot (C1) and accessibility tree (C2) conditions see a page genuinely built for people rather than a stripped-down agent shell. Sold-out sizes are visible but disabled on the page, so a pixel or tree agent can still attempt them. The structured conditions act through a JSON API. The screenshot (C1) and accessibility tree (C2) conditions drive the human pages in a real browser and never touch that API to act.

Why only the interface varies. If the store, the tasks, the grader, or the model changed at the same time as the interface, no difference could be attributed to the interface. So everything else is pinned.

Table I1. Held fixed versus varied.

Held fixed	Varied
MiniShop, one backend, one grader	How the agent <i>sees</i> the store
Ten frozen tasks	How the agent <i>acts</i> on it
One model per comparison; temperature 0; step cap 20	Presence of the view document, comparing view document (C4) against flat tools (C3)
Memory: task, prior actions, current observation	

One honest confound remains. The screenshot (C1) and accessibility tree (C2) act through coordinates and named elements, which cannot be identical to the structured conditions' function calls. Comparisons between human and structured surfaces therefore carry an action-vocabulary difference along with the representation difference. The comparison that isolates the *document* itself is flat tools (C3) against view document (C4).

Matched action grain. Both structured conditions expose the same seven operations:

`open_product`, `set_size`, `add_to_cart`, `go_catalog`, `go_checkout`, `set_address`, and `pay`. Flat tools (C3) stops there. The view document (C4) adds the current view, per-action `enabled` flags, and argument schemas. If the view document (C4) had fewer or simpler actions, any win could be explained as "fewer buttons." Matched grain removes that explanation. What matched grain does not separate is *why* the document helps when it helps: because it is compact, or because it carries constraints. Those two mechanisms stay entangled until the ablation described in Section 6.5 is run. Models are never mixed within a comparison.

The four conditions.

Condition	What the agent sees	How it acts
Screenshot (C1)	Rendered image of the human page	Click coordinates, type, scroll
Accessibility tree (C2)	Text tree of that page	Click or fill named elements
Flat tools (C3)	Function catalog, no current view	Call those functions
View document (C4)	JSON: view, state, entities, affordances	Invoke an enabled affordance with typed arguments

Tasks and grading. Ten frozen tasks: seven purchases that should succeed and three requests that should be refused. The grader inspects the backend order record and is never another model. A purchase counts only if the product, size, and shipping address all match. A refusal counts only if the forbidden order was never placed: t03 asks for a sold-out size, t05 asks to pay with an empty cart, t07 asks for another sold-out size. Because an agent that does nothing can technically pass a refusal, refusal results are reported separately and never inflate the headline success rate.

Table T1. The frozen task set.

ID	Expect	Instruction (abridged)
t01	success	Navy Crew Tee, size M, to 18 Cedar Ave, Portland
t02	success	Harbor Blue Tee, size L, to 18 Cedar Ave, Portland
t03	refusal	Harbor Blue Tee, size M (sold out)
t04	success	Studio Mug, to 9 Pine Street, Austin
t05	refusal	Pay with an empty cart
t06	success	Ash Hoodie, size S, to 4 Market Road, Seattle
t07	refusal	Signal Cap, size S (sold out)
t08	success	Field Notebook, to 100 King St, Boston
t09	success	Day Bottle and Crew Socks, to 18 Cedar Ave, Portland
t10	success	Navy Crew Tee, size L, to 77 Oak Lane, Denver

How a run works. One run is one combination of condition, task, and repeat. The agent loops (observe, act, observe) until the task completes or the twenty-step cap is hit. A full sweep is 4 conditions by 10 tasks by 5 repeats, 200 runs. Before any model runs, a deterministic oracle that replays the known-correct policy is executed on flat tools (C3) and view document (C4). It scores 100 percent on both, which establishes that any later miss belongs to the model, not to the harness

or the grader. Every run writes an independent trace with per-step usage, so the analysis reads finished artifacts rather than live state, and conditions are launched separately so one crash cannot take down the others.

4.1 A real view document (C4)

The document always has four parts. `view` names the screen the agent is on. `state` holds the session fields that matter on that screen. `entities` are the objects shown. `affordances` are the actions, each with an `enabled` flag and an `input` schema for its arguments.

The listing below is an actual document from MiniShop, not a sketch. The agent has opened the Navy Crew Tee, selected size L, and added it to the cart. It is still on the product page. Checkout has not been reached, so entering a shipping address is not yet a legal move, and the document says so: `set_address` and `pay` carry `enabled: false`, while `go_checkout` is available. The product-id list is shortened here for print; the live document lists every catalog id.

```
{
  "view": "product",
  "state": {
    "product_id": "tee-navy",
    "selected_size": "L",
    "cart_count": 1,
    "address": "",
    "order_count": 0
  },
  "entities": {
    "product": {
      "id": "tee-navy",
      "name": "Navy Crew Tee",
      "price": 28,
      "sizes_in_stock": ["S", "M", "L"]
    },
    "cart": [
      {
        "product_id": "tee-navy",
        "name": "Navy Crew Tee",
        "size": "L",
        "price": 28
      }
    ]
  },
  "affordances": [
    {
      "id": "open_product",
      "enabled": true,
      "input": {
        "type": "object",
        "properties": {
          "product_id": {
            "type": "string"
          }
        },
        "required": [
          "product_id"
        ]
      }
    },
    {
      "id": "set_size",
      "enabled": true,
      "input": {
        "type": "object",
        "properties": {
          "size": {
            "type": "string",
            "enum": [
              "S",
              "M",
              "L"
            ]
          }
        },
        "required": [
          "size"
        ]
      }
    },
    {
      "id": "add_to_cart",
      "enabled": true,
      "input": {
        "type": "object",
        "properties": {
          "product_id": {
            "type": "string"
          },
          "size": {
            "type": "string",
            "enum": [
              "S",
              "M",
              "L"
            ]
          }
        },
        "required": [
          "product_id",
          "size"
        ]
      }
    },
    {
      "id": "go_catalog",
      "enabled": true,
      "input": {
        "type": "object",
        "properties": {
          "product_id": {
            "type": "string"
          }
        },
        "required": [
          "product_id"
        ]
      }
    },
    {
      "id": "go_checkout",
      "enabled": true,
      "input": {
        "type": "object",
        "properties": {
          "product_id": {
            "type": "string"
          },
          "size": {
            "type": "string",
            "enum": [
              "S",
              "M",
              "L"
            ]
          }
        },
        "required": [
          "product_id",
          "size"
        ]
      }
    },
    {
      "id": "set_address",
      "enabled": false,
      "input": {
        "type": "object",
        "properties": {
          "address": {
            "type": "string"
          }
        },
        "required": [
          "address"
        ]
      }
    },
    {
      "id": "pay",
      "enabled": false,
      "input": {
        "type": "object",
        "properties": {
          "product_id": {
            "type": "string"
          },
          "size": {
            "type": "string",
            "enum": [
              "S",
              "M",
              "L"
            ]
          },
          "address": {
            "type": "string"
          }
        },
        "required": [
          "product_id",
          "size",
          "address"
        ]
      }
    }
  ]
}
```

Hold on to this example, because it is exactly the state where the one failure of the view document (C4) happens. The correct next move is `go_checkout`. In the model run, the model instead called `set_address` from this state, with `enabled: false` printed in front of it, and the backend rejected the call. The document told the truth and the model did not read it. Section 6.3 quantifies that failure.

Two properties make the document trustworthy as a measurement instrument. First, it is generated from backend state, so it is faithful by construction, and the enable rules mirror the ones the human page already uses: `add_to_cart` only when a valid size is selected, `set_address` only at checkout, `pay` only at checkout with a non-empty cart and address, and the size list always equal to what is actually in stock. Second, an automated consistency check forbids the document from ever advertising an action the backend would reject, so a good result for the view document (C4) cannot be an artifact of a surface that quietly disagrees with the application. Because it is written by hand against a known backend, it should be read as an upper bound on how good such a surface can be.

5. What we measure

Per run. Success (for refusal tasks, correctly declining). Input and output tokens summed over steps. Step count. Illegal actions (backend rejections). Malformed actions (replies that are not a usable action at all), counted separately so that formatting failures are never disguised as validity failures.

How tokens decompose. Each step costs a roughly constant overhead (instructions, task, prior actions) plus the current observation, whose size depends on the condition. A representation can therefore win in two ways: by being lighter per look, or by needing fewer looks.

Text conditions. For the accessibility tree (C2), flat tools (C3), and view document (C4), the observation is serialized text, and its cost is simply the length of that text under a fixed tokenizer. Two structural facts matter. Flat tools (C3) re-sends its full function schema with every call, which dominates its per-step cost. The view document (C4) pays for the document itself. Median per-step input along the shortest correct path:

Condition	System	Task and prior	Tool schema	Observation	Per-step total
Screenshot (C1), estimate	186	62	not sent	1105 (image)	1353
Accessibility tree (C2), measured	188	62	not sent	~650 (tree)	~900
Flat tools (C3)	159	62	490	25	736
View document (C4)	188	62	not sent	422	672

The flat tools (C3) schema (490 tokens) dwarfs its tiny status object (25 tokens). The view document (C4) spends its budget on the document itself (422 tokens). The accessibility tree (C2) is measured from live browser traces because it requires a rendered page. The other three rows are reconstructed deterministically with no model call.

The screenshot (C1) is priced differently. Images are billed by resolution tiles, not by tokenized content. Under the widely documented gpt-4o-style rule (fit within a 2048 by 2048 box, scale the short side toward 768 pixels, cover with 512-pixel tiles, charge a base plus a per-tile rate), our 1280 by 800 viewport costs six tiles, about 1105 image tokens, and a busier page at the same size costs no more. Providers also report image cost differently. OpenAI itemizes image tokens. Vertex and Gemini fold them into the total prompt count, which is why our screenshot (C1) runs report an itemized image count of zero alongside roughly 46,000 input tokens per task. The image is billed either way, so we read the screenshot (C1) cost from total input tokens. Tiling constants vary by provider, and we present these numbers as a labeled estimate rather than a universal formula.

Analysis. We report one table per model and condition, and a cost-reliability frontier of median input tokens against success rate. Because every condition runs the same tasks, we also compare the view document (C4) against flat tools (C3) *paired per task*: average each task over its repeats, take the difference, and bootstrap a 95 percent confidence interval over tasks, with success and refusal tasks reported separately. Models are never pooled.

Latency is not yet instrumented. Time weights round trips differently from tokens, so a timed re-run could shift the structured comparison without changing any token number. It is listed as future work.

6. Results

All results in this section come from one model on one store. The robust reading is the coarse ordering (human surfaces against structured ones). The fine margin between flat tools (C3) and view document (C4) is a provisional, single-model result.

6.1 What each representation costs before any model touches it

Replaying each task's shortest correct path and counting observation tokens with a fixed tokenizer gives a model-free baseline. One look costs a median of roughly 1105 tokens on the screenshot (C1), 512 on flat tools (C3) (the re-sent schema), and 453 on the view document (C4). Text is about 2.4 times cheaper per look than pixels, and the view document (C4) is slightly cheaper per look than flat tools (C3). This measures the observation term only. Real runs add step-count effects, and Section 6.3 shows they reverse the ordering of flat tools (C3) and view document (C4).

6.2 The apparatus, checked without a model

A deterministic oracle that replays the known-correct policy scores 100 percent on flat tools (C3) and view document (C4). Its token counts are locally tokenized rather than provider-billed, and the flat tools (C3) schema is not included in that count because the oracle loop delivers schemas out of band. The screenshot (C1) and accessibility tree (C2) are out of the oracle's scope, since they would need pixel-level or element-level oracles. The point of the control is attribution: when a model later misses, the miss belongs to the model.

6.3 The model run: **gemini-2.5-flash**, five repeats

One general model, temperature zero, step cap twenty, four conditions, ten tasks, five repeats: 200 runs. Confidence intervals are percentile bootstraps over the fifty task-repeat units in each cell.

Table 1. Per-condition results.

Condition	Success (95% CI)	Median input tokens (95% CI)	Median steps	Illegal/step	Malformed/step
Screenshot (C1)	30% [18, 42]	~46,338 [46,324, 46,368]	20	0.000	0.000
Accessibility tree (C2)	72% [60, 84]	~6,770 [5,835, 9,714]	7	0.198	0.116
Flat tools (C3)	100% [100, 100]	~3,134 [3,118, 3,154]	7	0.000	0.000
View document (C4)	92% [84, 98]	~4,455 [4,413, 4,491]	6	0.016	0.000

The whole sweep cost about 2.82 million input tokens and 35 thousand output tokens, about 93 cents at the recorded rates (0.30 dollars per million input, 2.50 per million output). The screenshot (C1) alone is about 73 percent of that bill (2.11 million input; the accessibility tree (C2) used 385 thousand, flat tools (C3) 138 thousand, view document (C4) 193 thousand). Screenshots dominate the cost before the question of whether they succeed is even asked.

Success. The screenshot (C1) passes 30 percent of runs and *none* of the seven purchase tasks; its passes are all refusals. The accessibility tree (C2) reaches 72 percent but earns it noisily, mixing backend rejections with malformed replies. Flat tools (C3) is perfect. The view document (C4) is near-perfect at 92 percent, and its entire deficit is one task, t10. With five repeats over ten tasks these intervals are wide, and the structured gap rests on that single task.

Tokens. The screenshot (C1) costs about ten times the view document (C4) and about fifteen times flat tools (C3) per task. Among the text conditions, flat tools (C3) is cheapest, the view document (C4) is next, and the accessibility tree (C2) is the most expensive and the most variable, because tree size tracks page content. Note the reversal against Section 6.1: per look, the view document (C4) is cheaper than flat tools (C3), but in real runs the document grows as the cart fills, and flat tools (C3)'s tiny status objects win on total cost.

Steps. The view document (C4) finishes fastest, median six steps, against seven for flat tools (C3) and the accessibility tree (C2). The screenshot (C1) sits at the cap of twenty because it never completes a purchase and simply spends its budget.

Errors. The screenshot (C1) and flat tools (C3) commit zero illegal and zero malformed actions. The view document (C4) commits 0.016 illegal actions per step, all on t10. The accessibility tree (C2) is the outlier, 0.198 illegal and 0.116 malformed per step: the human tree invites named clicks the backend rejects, and the model periodically produces replies the loop cannot parse at all. One caution: the screenshot (C1)'s clean error rate is not a virtue. A coordinate click rarely trips the backend's validity check, so its failures show up as capped steps and spent tokens instead.

Table 2. Per-task readout. Each cell is success rate, median steps, and median input tokens over five repeats. Refusal tasks are marked *r*.

Task		Screenshot (C1)	Accessibility tree (C2)	Flat tools (C3)	View document (C4)
t01		0% · 20 · 46,368	20% · 12 · 11,702	100% · 7 · 3,154	100% · 6 · 4,498
t02		0% · 20 · 46,368	100% · 11 · 10,809	100% · 7 · 3,154	100% · 6 · 4,486
t03	<i>r</i>	100% · 2 · 4,184	100% · 2 · 1,997	100% · 2 · 1,081	100% · 2 · 1,597
t04		0% · 20 · 46,268	0% · 6 · 5,835	100% · 7 · 3,118	100% · 6 · 4,423
t05	<i>r</i>	100% · 20 · 46,018	100% · 3 · 2,605	100% · 2 · 713	100% · 2 · 1,528
t06		0% · 20 · 46,328	100% · 10 · 9,714	100% · 7 · 3,144	100% · 6 · 4,491
t07	<i>r</i>	100% · 20 · 46,348	100% · 2 · 1,990	100% · 2 · 1,079	100% · 1 · 905
t08		0% · 20 · 46,324	100% · 7 · 6,690	100% · 7 · 3,124	100% · 6 · 4,413
t09		0% · 20 · 46,428	0% · 20 · 21,658	100% · 12 · 5,833	100% · 10 · 7,867
t10		0% · 20 · 46,368	100% · 7 · 6,770	100% · 7 · 3,154	20% · 6 · 4,544

The table makes four patterns visible. The screenshot (C1) hits the step cap on every task it does not refuse, at a nearly flat 46 thousand tokens, because image cost is set by the viewport. The accessibility tree (C2) is erratic, failing t01 and t04 outright and capping out on the two-item task t09 with the heaviest text-condition token bill (about 21.7 thousand). Flat tools (C3) is uniformly clean. The view document (C4) matches it everywhere except t10, usually one step faster and always at more token cost.

The paired comparison. Averaging each task over its five repeats and taking the difference (view document (C4) minus flat tools (C3)) gives a per-task comparison with a bootstrap interval over tasks.

Table 3. Paired per-task differences. Positive means the view document (C4) is higher.

Task set	n	Metric	Mean difference	95% CI	Excludes 0?
success	7	success	-0.114	[-0.343, 0.000]	no
success	7	input tokens	+1,404	[+1,317, +1,549]	yes
success	7	steps	-1.17	[-1.51, -1.00]	yes
refusal	3	success	0.000	[0.000, 0.000]	no
refusal	3	input tokens	+386	[-174, +815]	no
refusal	3	steps	-0.33	[-1.00, 0.00]	no

On purchase tasks, the view document (C4) saves steps (the interval excludes zero, which is precisely what a current-view document is for) and costs more tokens (the interval excludes zero, and all seven tasks point the same way). It does not measurably change success (the interval includes zero, and the entire point estimate is t10). So on this store and this model, flat tools (C3) is preferred on the frontier: cheaper, and at least as reliable. The strong form of H2, that the view document (C4) dominates, is not supported here. On refusal tasks nothing separates the two.

The one failure, told plainly. Task t10 asks the agent to buy a Navy Crew Tee in size L and ship it to 77 Oak Lane, Denver. The legal order of operations is: open the product, add it to the cart, go to checkout, and only then set the address. The model instead tried to set the address while still on the product page. That is exactly the state printed in Section 4.1, where `set_address` is marked `enabled: false`. The backend rejected the call, the run burned the steps, and the task failed. The document was correct, the backend agreed with it, and an automated check pins that agreement, so this is not an interface bug. It is a model ignoring a rule the interface stated. The failure reproduced on four of five sweep repeats and five of five in a follow-up diagnostic, always as the same premature call. The lesson generalizes: a truthful document cannot help a model that does not read `enabled`. How often a model attempts actions the interface marked unavailable (its ignored-affordance rate) is a quantity the next model must measure.

6.4 A second model (not yet run)

Whether flat tools (C3) stays preferred is an open question about models, not about the store. The single most informative datum from a second general model is whether it respects `enabled: false` on t10.

6.5 An ablation (not yet run)

Stripping the `enabled` flags and argument constraints out of the view document (C4), while keeping the view and entities, would separate the document's two entangled mechanisms: compactness and constraint-carrying. If the step savings vanish, the constraints were doing the work.

7. Discussion

7.1 What this cell establishes

- **Representation moves every measured quantity (RQ1).** Success from 30 to 100 percent, median tokens by a factor of about fifteen, steps from six to the cap, and error rates from zero to 0.198 illegal per step, all from re-describing the same store to the same model.
- **Human surfaces are dominated here (RQ2, first half).** Both the screenshot (C1) and the accessibility tree (C2) are more expensive and less reliable than either structured condition.
- **The document does not yet beat the catalog (RQ2, second half).** Flat tools (C3) is preferred over the view document (C4) on tokens and success, while the view document (C4) reliably saves steps. The mechanism a purpose-built surface promises is visible, and it is not yet large enough to pay its token bill on this store.
- **The one failure is behavioral, not structural.** On t10 the document told the truth and the model ignored it. That distinction only exists because the surface is checked against the backend, which is an argument for building agent surfaces that can be checked.

It is worth saying why this store is the hard case for the view document (C4). MiniShop is tiny: few views, short tasks, little state to track. That is exactly where flat tools (C3)'s weakness, making the model carry the state in its head, costs the least. A representation that pays per-step to describe state earns the least where there is the least state to describe.

7.2 What later cells must show

These are predictions, not results.

1. **A second general model (RQ3).** If it obeys `enabled: false` on t10, the view document (C4) closes the success gap. If it also ignores the flag, the miss is a property of current models rather than of one model. Either answer is informative. Models stay unpooled, and if a model lacks vision, the screenshot (C1) is dropped and reported as such.
2. **The ablation.** Decides whether the document's value is compactness or constraints.
3. **Latency.** The view document (C4) uses fewer round trips. Priced in time instead of tokens, the verdict could flip without any number in Table 1 changing.
4. **A harder application.** More views, more state, more ways to act illegally. That is where the state-tracking tax on flat tools (C3) is hypothesized to grow, and where the trade should tilt. Until that cell exists, "flat tools (C3) wins on MiniShop" must not be read as "do not build agent surfaces."

8. How an application could ship a view document (C4)

The measured document is hand-written, an upper bound. The encouraging fact for adoption is that almost everything in it is validity logic the application already has. A human page already disables

Add to cart until a size is chosen, already validates the address field, already hides Pay until the cart and address are set. Shipping a view document (C4) is mostly *serializing predicates that already exist*, not authoring new business rules. The backend remains the authority, and the document must never advertise an action the backend would reject, which is a testable property.

Five ways to produce the document, from highest leverage to most general:

1. **Framework-emitted.** The developer declares views, state, and actions once, and the framework renders both the human page and the view document (C4) from that single declaration. Near-zero marginal effort and no drift between the two surfaces, but it requires framework support that is not yet common.
2. **Projected from a spec.** Applications with a GraphQL schema, form validation schemas, or server-driven UI already hold their capabilities in machine-readable form. Types become entities, mutations become affordances, validation rules become `enabled` logic. Fidelity is high wherever the spec is the real authority.
3. **Hand-written per view (this study).** Faithful by construction and pinned by tests, but bespoke per view, so it suits a few high-value flows rather than an ecosystem. As a size reference, MiniShop's document generator is roughly three to four times the code of its flat tools (C3) catalog, and it mirrors checks the human page already makes.
4. **Compiled from the human page.** A compiler reads the DOM, ARIA roles, and disabled attributes and emits a document. It costs the application nothing, but it inherits only what the page happens to render, and it can disagree with the backend, so it must be validated before it is trusted.
5. **Model-extracted.** A model reads the screen and writes the document. It works on any application with zero cooperation, and it inverts the economics: the cost returns on every step, along with the risk of hallucinated actions. A fallback, not the goal.

The underlying trade. A view document (C4) pays once, at authoring or compile time, for the ability to state what is true and allowed at every step. Flat tools (C3) pays forever at runtime, because the model must re-derive state from its own history on every step, for the life of the deployment. On a tiny store that runtime tax is small, which is exactly what Section 6 measures. It grows with views, state, and ways to be wrong, and it is amortized across every agent that ever visits the application.

A short recipe for a developer today. Enumerate the views. For each view, expose the state fields, the entities, and the actions, reusing the exact enable-conditions and validators the human page already has. Serve the document alongside the page. Keep backend enforcement in place. Add one automated check that the document never advertises what the backend would reject.

Table 4. The five paths at a glance.

Path	Application effort	Coverage	Fidelity	Best fit
Framework-emitted	Low at scale	High	High, no drift	New applications on agent-aware frameworks
Projected from a spec	Low to medium	Wherever the spec rules	High there	GraphQL, form schemas, server-driven UI
Hand-written per view	Medium per app	What is authored	High, test-pinned	A few high-value flows; this paper
Compiled from the human page	Near zero	Any app with a page	Lossy, must be checked	Legacy applications
Model-extracted	Near zero	Universal	Lowest, cost recurs	Fallback only

Measuring the two generated paths (compiled and model-extracted) against the hand-written upper bound established here is future work.

9. Limitations

Scope. One model, one small synthetic store, ten tasks. The paired intervals resample over seven purchase tasks and three refusals, so they are wide by construction. Nothing here generalizes beyond a controlled shopping task yet.

The document is an upper bound. A hand-written view document (C4) is as good as such a surface gets. A sloppier generated one would do worse, and by how much is unmeasured.

Entangled mechanisms. Compactness and constraint-carrying are not separated until the ablation runs.

Refusals can pass by inaction. They are therefore reported apart from purchases throughout.

Image accounting is provider-specific. Some providers itemize image tokens and some fold them into the prompt count, so the screenshot (C1) is read from total input tokens, and the tiling numbers are a labeled estimate.

Temperature zero is not determinism. Results are averaged over five repeats. Flat tools (C3) was stable across repeats, the view document (C4) varied only on t10, and the accessibility tree (C2) was noisy throughout.

The memory policy is fixed. Agents carry only the task, their prior actions, and the current observation. This is the same for every condition, but it also means flat tools (C3) cannot compensate by accumulating a long transcript, and other memory policies are untested.

The screenshot (C1) error rate understates its errors. Coordinate clicks rarely trigger the backend's validity check, so for the screenshot (C1) the informative numbers are success, steps, and tokens.

Wide intervals are not equality. An interval that includes zero, as the paired success difference does, means the data cannot distinguish the conditions at this sample size, not that they are the same.

Specialized models are out of scope. The main line uses general models that can run all four conditions through one API. Pixel-native computer-use models are a planned final check, and the second general model has not yet been run.

10. Future work

Nearest first: instrument latency and re-run timed; run a second and then a third general model, reporting each separately; run the constraint-stripping ablation; report an ignored-affordance rate per model. After that: a second, more complex application, which is the main lever on external validity; a validated Agent Experience metric suite, which is a separate paper; and automatic generation of the agent surface, compiled from the page or extracted by a model, measured against the hand-written upper bound established here.

Appendix: reproducing the numbers

The store, the four-condition harness, the grader, and every trace behind the tables live in the project repository, with no hidden state. The sequence is fixed. A scripted non-model policy must first complete every purchase task on flat tools (C3) and view document (C4) and be refused on the refusals. The model-free baseline of Section 6.1 then runs with a fixed tokenizer and no API call. The oracle of Section 6.2 validates the pipeline. The model sweep writes its 200 traces, and the report layer aggregates them into Tables 1 through 3 with their bootstrap intervals. Dependencies and the browser used by the screenshot (C1) and accessibility tree (C2) conditions are pinned in the repository, and this manuscript rebuilds to HTML and PDF from its Markdown source.

References

Verified sources only. Each entry was checked against the primary paper or official documentation. No other citations are asserted.

- WebArena: A Realistic Web Environment for Building Autonomous Agents.
<https://arxiv.org/abs/2307.13854>
- VisualWebArena: Evaluating Multimodal Agents on Realistic Visually Grounded Web Tasks. <https://aclanthology.org/2024.acl-long.50/>

- Mind2Web: Towards a Generalist Agent for the Web. <https://arxiv.org/abs/2306.06070>
- WebShop: Towards Scalable Real-World Web Interaction with Grounded Language Agents. https://proceedings.neurips.cc/paper_files/paper/2022/file/82ad13ec01f9fe44c01cb91814fd7b8c-Paper-Conference.pdf
- AndroidWorld: A Dynamic Benchmarking Environment for Autonomous Agents. <https://arxiv.org/abs/2405.14573>
- MiniWoB++ (Reinforcement Learning on Web Interfaces using Workflow-Guided Exploration). <https://arxiv.org/abs/1802.08802> ; docs: <https://miniwob.farama.org/>
- Anthropic: Introducing computer use (Claude 3.5 Sonnet). <https://www.anthropic.com/news/3-5-models-and-computer-use>
- UI-TARS: Pioneering Automated GUI Interaction with Native Agents. <https://arxiv.org/abs/2501.12326>
- SeeClick: Harnessing GUI Grounding for Advanced Visual GUI Agents. <https://aclanthology.org/2024.acl-long.505/>
- UIFormer / From User Interface to Agent Interface: Efficiency Optimization of UI Representations for LLM Agents. <https://arxiv.org/abs/2512.13438>
- Prune4Web: DOM Tree Pruning Programming for Web Agent. <https://ojs.aaai.org/index.php/AAAI/article/download/40772/44733>
- A2UI (Agent-to-UI) Protocol. <https://a2ui.org/>
- Model Context Protocol (MCP). <https://modelcontextprotocol.io/>